



Experiment No. 5
Apply appropriate Unsupervised Learning Technique on the Wholesale Customers Dataset
Date of Performance: 21 / 08 /2023
Date of Submission: 03 / 09 /2023



Aim: Apply appropriate Unsupervised Learning Technique on the Wholesale Customers Dataset.


Objective: Able to perform various feature engineering tasks, apply Clustering Algorithm on the given dataset.

Theory:

It is basically a type of unsupervised learning method. An unsupervised learning method is a method in which we draw references from datasets consisting of input data without labeled responses. Generally, it is used as a process to find meaningful structure, explanatory underlying processes, generative features, and groupings inherent in a set of examples.

Clustering is the task of dividing the population or data points into a number of groups such that data points in the same groups are more similar to other data points in the same group and dissimilar to the data points in other groups. It is basically a collection of objects on the basis of similarity and dissimilarity between them.

For example: The data points in the graph below clustered together can be classified into one single group. We can distinguish the clusters, and we can identify that there are 3 clusters in the below picture.





Dataset:

This data set refers to clients of a wholesale distributor. It includes the annual spending in monetary units (m.u.) on diverse product categories. The wholesale distributor operating in different regions of Portugal has information on annual spending of several items in their stores across different regions and channels. The dataset consist of 440 large retailers annual spending on 6 different varieties of product in 3 different regions (lisbon , oporto, other) and across different sales channel (Hotel, channel)

Detailed overview of dataset

Records in the dataset = 440 ROWS

Columns in the dataset = 8 COLUMNS

FRESH: annual spending (m.u.) on fresh products (Continuous)

MILK:- annual spending (m.u.) on milk products (Continuous)

GROCERY:- annual spending (m.u.) on grocery products (Continuous)

FROZEN:- annual spending (m.u.) on frozen products (Continuous)

DETERGENTS_PAPER :- annual spending (m.u.) on detergents and paper products (Continuous)

DELICATESSEN:- annual spending (m.u.)on and delicatessen products (Continuous);

CHANNEL: - sales channel Hotel and Retailer

REGION:- three regions (Lisbon, Oporto, Other)

Code:



Conclusion :-

Using Clustered Data: Data clustering is vital for identifying distinct customer groups, enabling tailored strategies, optimizing operations, and boosting business success.

Cluster 0: "Diverse Shoppers" - Moderate purchases, balanced marketing.

Cluster 1: "Freshness Enthusiasts" - High demand for fresh, fast delivery.

Cluster 2: "Budget-Conscious Buyers" - Smaller purchases, cost-effective options.

Cluster 3: "High-Volume Demands" - Premium and efficient delivery for high-volume customers.

Adapting Delivery Schemes: Aligning delivery with customer preferences enhances satisfaction and business growth.

Cluster 0: Cost-effective, reliable options.

Cluster 1: Rapid, temperature-controlled delivery.

Cluster 2: Consolidated or slower delivery for cost reduction.



Cluster 3: Premium, bulk delivery for high-volume needs.

Importing Libs

```
# https://www.kaggle.com/code/ahmedhisham73/kmeans
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv('/content/Wholesale.csv')
df.head(10)
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen	
0	2	3	12669	9656	7561	214	2674	1338	
1	2	3	7057	9810	9568	1762	3293	1776	
2	2	3	6353	8808	7684	2405	3516	7844	
3	1	3	13265	1196	4221	6404	507	1788	
4	2	3	22615	5410	7198	3915	1777	5185	
5	2	3	9413	8259	5126	666	1795	1451	
6	2	3	12126	3199	6975	480	3140	545	
7	2	3	7579	4956	9426	1669	3321	2566	
8	1	3	5963	3648	6192	425	1716	750	
9	2	3	6006	11093	18881	1159	7425	2098	

Data exploration:

```
print("Column names:")
print(df.columns)

Column names:
Index(['Channel', 'Region', 'Fresh', 'Milk', 'Grocery', 'Frozen',
      'Detergents_Paper', 'Delicassen'],
      dtype='object')
```

```
print("Data types:")
print(df.dtypes)

Data types:
Channel          int64
Region          int64
Fresh           int64
Milk            int64
Grocery         int64
Frozen          int64
Detergents_Paper int64
Delicassen      int64
dtype: object
```

```
print("Missing values per column:")
print(df.isnull().sum())

Missing values per column:
Channel      0
Region      0
Fresh        0
Milk         0
Grocery      0
Frozen       0
Detergents_Paper 0
Delicassen   0
dtype: int64
```

```
print("Descriptive Statistics:")
print(df.describe())

print("Number of duplicate rows: ", df.duplicated().sum())

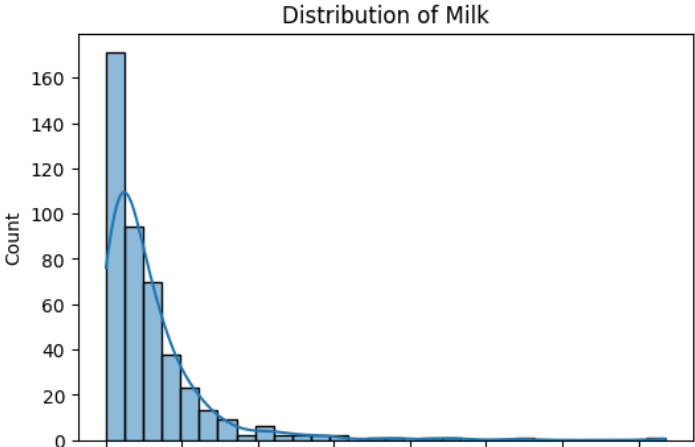



for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.histplot(df[column], bins=30, kde=True)
    plt.title(f'Distribution of {column}')
    plt.show()
```





Descriptive Statistics:

	Channel	Region	Fresh	Milk	Grocery \
count	440.000000	440.000000	440.000000	440.000000	440.000000
mean	1.322727	2.543182	12000.297727	5796.265909	7951.277273
std	0.468052	0.774272	12647.328865	7380.377175	9503.162829
min	1.000000	1.000000	3.000000	55.000000	3.000000
25%	1.000000	2.000000	3127.750000	1533.000000	2153.000000
50%	1.000000	3.000000	8504.000000	3627.000000	4755.500000
75%	2.000000	3.000000	16933.750000	7190.250000	10655.750000
max	2.000000	3.000000	112151.000000	73498.000000	92780.000000


	Frozen	Detergents_Paper	Delicassen
count	440.000000	440.000000	440.000000
mean	3071.931818	2881.493182	1524.870455
std	4854.673333	4767.854448	2820.105937
min	25.000000	3.000000	3.000000
25%	742.250000	256.750000	408.250000
50%	1526.000000	816.500000	965.500000
75%	3554.250000	3922.000000	1820.250000
max	60869.000000	40827.000000	47943.000000

Number of duplicate rows: 0





```
# Heatmap for correlation between variables
plt.figure(figsize=(6, 4))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()
```

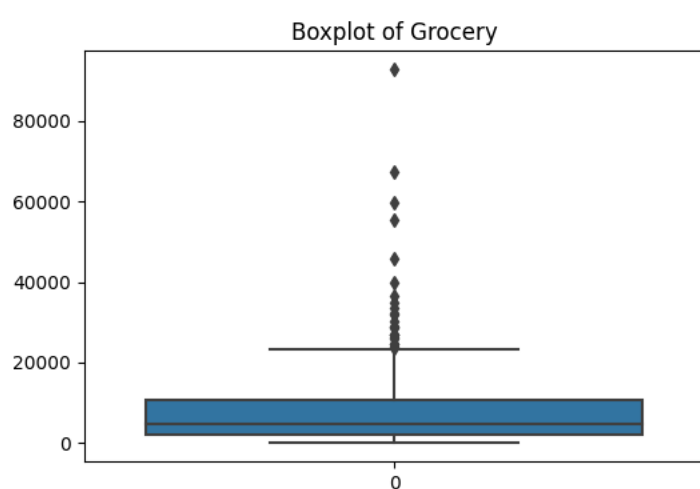

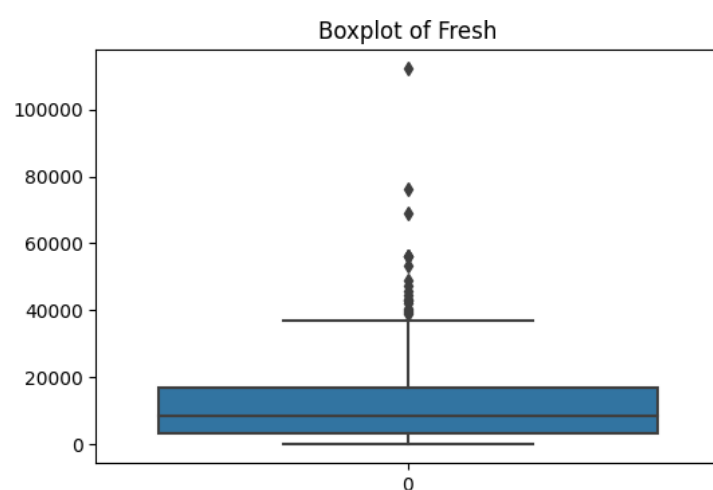


Checking for Outliers

```
# Draw boxplots for all features
for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.boxplot(df[column])
    plt.title(f'Boxplot of {column}')
    plt.show()

# Function to detect outliers
def detect_outliers(dataframe, column):
    Q1 = dataframe[column].quantile(0.25)
    Q3 = dataframe[column].quantile(0.75)
    IQR = Q3 - Q1
    outliers = dataframe[(dataframe[column] < Q1 - 1.5*IQR) | (dataframe[column] > Q3 + 1.5*IQR)]
    return outliers

# Detect and print number of outliers for each feature
for column in df.columns:
    outliers = detect_outliers(df, column)
    print(f'Number of outliers in {column}: {len(outliers)}')
```





Boxplot of Frozen

Elbow Method to identify number of cluster Required

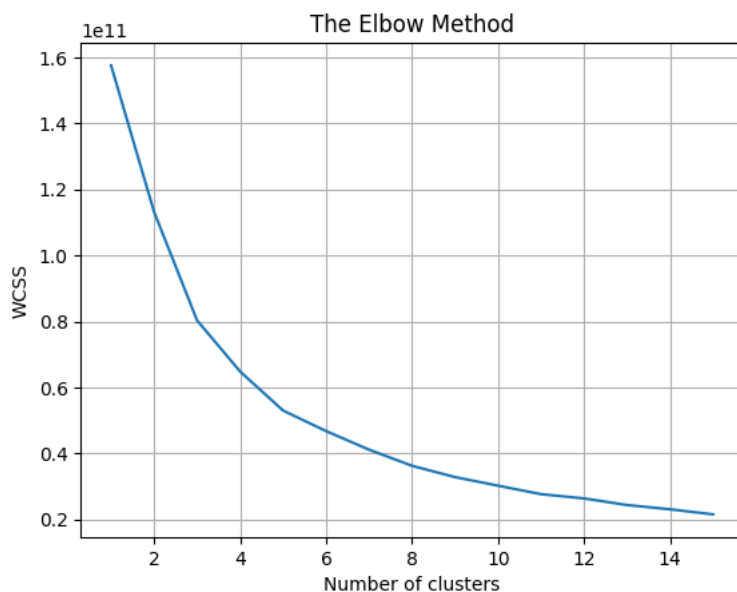
```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)

from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Calculate WCSS for different number of clusters
wcss = []
max_clusters = 15
for i in range(1, max_clusters+1):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(df)
    wcss.append(kmeans.inertia_)

# Plot the WCSS values
plt.plot(range(1, max_clusters+1), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()
```

[illegible]

KMeans

```
from sklearn.cluster import KMeans

# Build the model
kmeans = KMeans(n_clusters=4, init='k-means++', random_state=42)
kmeans.fit(df)

# Get cluster labels
cluster_labels = kmeans.labels_

# Add cluster labels to your original dataframe
df['Cluster'] = cluster_labels

print(df.head())
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	\
0	2	3	12669	9656	7561	214		2674
1	2	3	7057	9810	9568	1762		3293
2	2	3	6353	8808	7684	2405		3516
3	1	3	13265	1196	4221	6404		507
4	2	3	22615	5410	7198	3915		1777

	Delicassen	Cluster
0	1338	2
1	1776	0
2	7844	2
3	1788	2
4	5185	1

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 1 to 10 in version 0.25. For now the warning is only displayed for non-optimal values of n_init. Please set n_init to the desired value. warnings.warn(

```
# Add cluster labels to the DataFrame
df['Cluster'] = kmeans.labels_

# Check the size of each cluster
print("Cluster Sizes:\n", df['Cluster'].value_counts())

# Check the characteristics of each cluster
for i in range(4):
    print("\nCluster ", i)
    print(df[df['Cluster'] == i].describe())
```

std	0.347839	0.773111	15406.720722	5059.930789	4289.538677
min	1.000000	1.000000	22096.000000	286.000000	471.000000
25%	1.000000	3.000000	26434.750000	2158.750000	2618.750000
50%	1.000000	3.000000	30818.000000	3954.500000	5058.500000
75%	1.000000	3.000000	40604.250000	7103.500000	8219.000000
max	2.000000	3.000000	112151.000000	29627.000000	18148.000000

	Frozen	Detergents_Paper	Delicassen	Cluster
count	58.000000	58.000000	58.000000	58.0
mean	6298.655172	1064.000000	2316.724138	1.0
std	8840.373423	1317.904703	2409.193705	0.0
min	127.000000	10.000000	3.000000	1.0
25%	1370.750000	284.000000	975.500000	1.0
50%	3662.000000	561.500000	1535.500000	1.0
75%	8674.000000	1135.500000	2798.000000	1.0
max	60869.000000	5058.000000	14351.000000	1.0

Cluster	2				
	Channel	Region	Fresh	Milk	Grocery \
count	276.000000	276.000000	276.000000	276.000000	276.000000
mean	1.152174	2.536232	9087.463768	3027.427536	3753.514493
std	0.359842	0.778431	6218.787958	2599.933332	2716.555045
min	1.000000	1.000000	3.000000	55.000000	3.000000
25%	1.000000	2.000000	3454.250000	1133.750000	1755.500000
50%	1.000000	3.000000	8257.500000	2193.000000	2849.000000
75%	1.000000	3.000000	13582.750000	4192.500000	5231.000000
max	2.000000	3.000000	23257.000000	18664.000000	13462.000000

	Frozen	Detergents_Paper	Delicassen	Cluster
count	276.000000	276.000000	276.000000	276.0
mean	2817.985507	1003.003623	1040.525362	2.0
std	3614.905029	1233.205498	1013.744595	0.0
min	47.000000	3.000000	3.000000	2.0
25%	779.000000	199.750000	360.750000	2.0
50%	1571.000000	436.000000	713.500000	2.0
75%	3505.250000	1322.750000	1417.250000	2.0
max	35009.000000	5316.000000	7844.000000	2.0





Cluster 3						
	Channel	Region	Fresh	Milk	Grocery	\
count	11.000000	11.000000	11.000000	11.000000	11.000000	
mean	1.909091	2.545455	19888.272727	36142.363636	45517.454545	
std	0.301511	0.820200	14488.239473	18433.364784	21771.475666	
min	1.000000	1.000000	85.000000	4980.000000	20170.000000	
25%	2.000000	2.500000	8881.500000	25302.500000	32074.000000	
50%	2.000000	3.000000	16117.000000	36423.000000	39694.000000	
75%	2.000000	3.000000	31157.500000	45073.500000	57584.500000	
max	2.000000	3.000000	44466.000000	73498.000000	92780.000000	





	Frozen	Detergents_Paper	Delicassen	Cluster
count	11.000000	11.000000	11.000000	11.0
mean	6328.909091	21417.090909	8414.000000	3.0
std	10355.038662	12078.543310	13829.020486	0.0
min	36.000000	239.000000	903.000000	3.0
25%	1006.500000	18750.000000	1720.000000	3.0
50%	3254.000000	20070.000000	2944.000000	3.0
75%	5950.000000	25466.000000	5797.500000	3.0
max	36534.000000	40827.000000	47943.000000	3.0

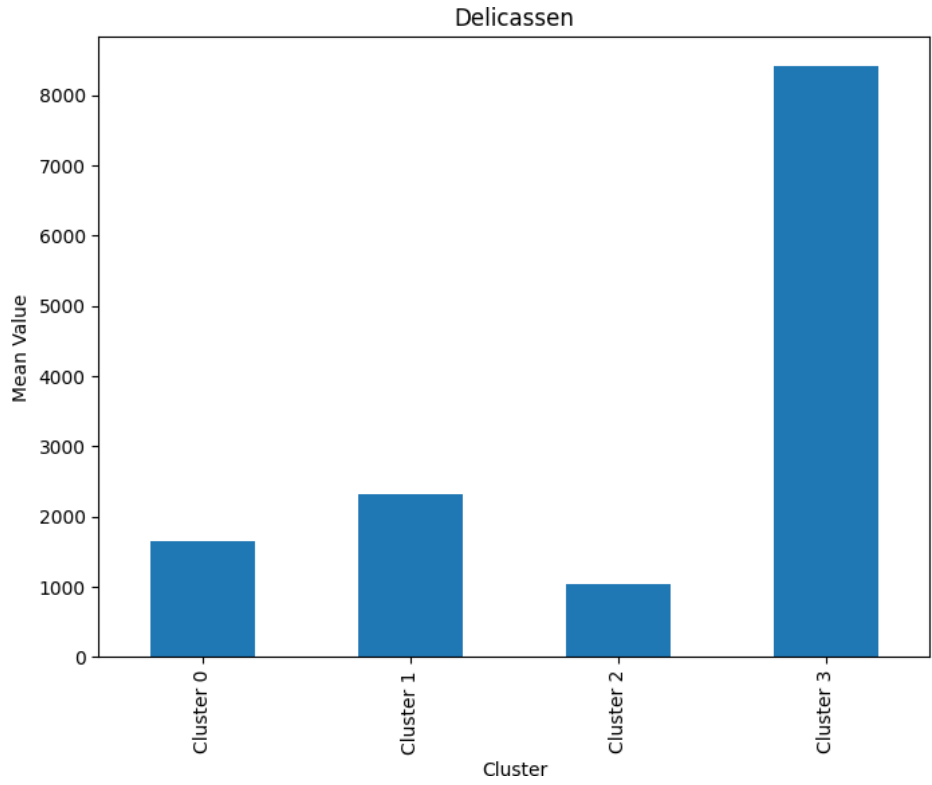

```
# Calculate the mean values for each feature per cluster
cluster_means = df.groupby('Cluster').mean()

# Transpose the DataFrame so that the features are the rows (this will make plotting easier)
cluster_means = cluster_means.transpose()

# Create bar plot for each feature
for feature in cluster_means.index:
    cluster_means.loc[feature].plot(kind='bar', figsize=(8,6))
    plt.title(feature)
    plt.ylabel('Mean Value')
    plt.xticks(ticks=range(4), labels=['Cluster 0', 'Cluster 1', 'Cluster 2', 'Cluster 3'])
    plt.show()
```








```
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

# Apply PCA and fit the features selected
pca = PCA(n_components=2)
principalComponents = pca.fit_transform(df.drop('Cluster', axis=1))

# Create a DataFrame with the two components
PCA_components = pd.DataFrame(principalComponents, columns=['Principal Component 1', 'Principal Component 2'])

# Concatenate the clusters labels to the DataFrame
PCA_components['Cluster'] = df['Cluster']

# Plot the clustered dataset
plt.figure(figsize=(8,6))
plt.scatter(PCA_components['Principal Component 1'], PCA_components['Principal Component 2'], c=PCA_components['Cluster'])
plt.title('Clusters in PCA 2D Space')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster')
plt.show()
```



● x

