



Experiment No. 6
Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance: 04/09/2023
Date of Submission: 13/09/2023



Aim: Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—create an ensemble of classifiers. Each one gives a weighted vote.

Input:

- D , a set of d class labelled training tuples
- k , the number of rounds (one classifier is generated per round)
- a classification learning scheme

Output: A composite model

Method

1. Initialize the weight of each tuple in D is $1/d$
2. For $i=1$ to k do // for each round
3. Sample D with replacement according to the tuple weights to obtain D_i
4. Use training set D_i to derive a model M_i
5. Compute $\text{error}(M_i)$, the error rate of M_i
6. $\text{Error}(M_i) = \sum_j w_j * \text{err}(X_j)$
7. If $\text{Error}(M_i) > 0.5$ then
8. Go back to step 3 and try again
9. endif
10. for each tuple in D_i that was correctly classified do
11. Multiply the weight of the tuple by $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for



To use the ensemble to classify tuple X

1. Initialize the weight of each class to 0
2. for $i=1$ to k do // for each classifier
3. $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifiers vote
4. $C = M_i(X)$ // get class prediction for X from M_i
5. Add w_i to weight for class C
6. end for
7. Return the class with the largest weight.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.



capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code



Conclusion:

1. Accuracy, confusion matrix, precision, recall and F1 score obtained.

The Adaboost model exhibited an overall accuracy of 86.37% in predicting income levels. In the specific category of income " $= \$50K$," the model demonstrated commendable precision at 0.88, indicating a high accuracy of positive predictions. However, for the category " $> \$50K$," the precision was slightly lower at 0.79, suggesting room for enhancement in correctly identifying instances within this income bracket. The model's recall, a measure of its ability to capture all relevant instances, was notably high at 0.94 for " $= \$50K$," signifying effective identification of individuals with this income level. On the other hand, the recall for " $> \$50K$ " was 0.63, indicating a moderate performance in correctly identifying individuals in this income category. The F1 score, which balances precision and recall, was 0.91 for " $= \$50K$," suggesting a harmonious blend of accuracy and completeness in predictions within this income range. For " $> \$50K$," the F1 score was 0.70, indicating a reasonable balance but with room for improvement.

confusion matrix				
[[6374 379]				
[854 1444]]				
	precision	recall	f1-score	support
0	0.88	0.94	0.91	6753
1	0.79	0.63	0.70	2298
accuracy			0.86	9051
macro avg	0.84	0.79	0.81	9051
weighted avg	0.86	0.86	0.86	9051

2 Comparison of Boosting and Random Forest Algorithms:

The boosting algorithms—AdaBoost, Gradient Boosting, and XGBoost—demonstrate superior performance in terms of accuracy, precision, and F1 score compared to the Random Forest Classifier. While the Random Forest Classifier achieves a commendable accuracy of around 86.37% and maintains a balanced F1 score, it falls slightly behind the boosting algorithms in terms of overall predictive metrics. The boosting algorithms exhibit higher precision and recall for the positive class, reflecting their enhanced capability to accurately identify and categorize individuals with high incomes. This suggests that the boosting algorithms excel in pinpointing cases related to high-income individuals, showcasing a notable advantage over the Random Forest Classifier in this regard. Across the Adult Census Income Dataset, all models perform effectively, meeting expectations for accurate income level classification. Gradient Boosting, in particular, stands out with a slightly elevated accuracy and F1 score compared to the Random Forest Classifier. This implies that Gradient Boosting may provide a more refined and accurate prediction on this specific dataset, reinforcing its potential as a robust model for income classification.

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import io
from sklearn.metrics import accuracy_score, precision_score, f1_score, \
    confusion_matrix, classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
# all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
[ ]: file = ('/content/adult.csv')
df = pd.read_csv(file)
```

```
[ ]: print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
0	90	?	77053	HS-grad	9	Widowed	
1	82	Private	132870	HS-grad	9	Widowed	
2	66	?	186061	Some-college	10	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	

		relationship	race	sex	capital.gain	\
0		? Not-in-family	White	Female	0	
1	Exec-managerial	Not-in-family	White	Female	0	
2		? Unmarried	Black	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	

	cap	ital.loss	hours.per.week	native.country	income
1	4356	4356	40	United-States	<=50K
2	4356	4356	18	United-States	<=50K
3	4356	4356	40	United-States	<=50K
4	3900	3900	40	United-States	<=50K
5	3900	3900	40	United-States	<=50K

```
[ ]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
5   marital.status         32561 non-null  object
6   occupation             32561 non-null  object
7   relationship           32561 non-null  object
8   race                   32561 non-null  object
9   sex                    32561 non-null  object
10  capital.gain           32561 non-null  int64
11  capital.loss           32561 non-null  int64
12  hours.per.week         32561 non-null  int64
13  native.country         32561 non-null  object
14  income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
None
```

```
[ ]: #Count the occuring of the '?' in all the columns
for i in df.columns:
    t = df[i].value_counts()
    index = list(t.index)
    print ("Count of ? in", i, end=" ")
    for i in index:
        temp = 0
        if i == '?':
            print (t['?'])
            temp = 1
            break
    if temp == 0:
        print ("0")
```

Count of ? in age 0

Count of ? in workclass 1836
 Count of ? in fnlwgt 0
 Count of ? in education 0
 Count of ? in education.num 0
 Count of ? in marital.status 0
 Count of ? in occupation 1843
 Count of ? in relationship 0
 Count of ? in race 0
 Count of ? in sex 0
 Count of ? in capital.gain 0
 Count of ? in capital.loss 0
 Count of ? in hours.per.week 0
 Count of ? in native.country 583
 Count of ? in income 0

```
[ ]: df=df.loc[(df['workclass'] != '?') & (df['native.country'] != '?')]
print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K
5	3770	45	United-States	<=50K
6	3770	40	United-States	<=50K

```
[ ]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	0
4	3900	40	United-States	0
5	3900	40	United-States	0
6	3770	45	United-States	0
7	3770	40	United-States	0

```
[ ]: df_more=df.loc[df["income"] == 1]
print(df_more.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
7	74	State-gov	88638	Doctorate	16	Never-married	
10	45	Private	172274	Doctorate	16	Divorced	
11	38	Self-emp-not-inc	164526	Prof-school	15	Never-married	
12	52	Private	129177	Bachelors	13	Widowed	
13	32	Private	136204	Masters	14	Separated	

	occupation	relationship	race	sex	capital.gain	\
7	Prof-specialty	Other-relative	White	Female	0	
10	Prof-specialty	Unmarried	Black	Female	0	
11	Prof-specialty	Not-in-family	White	Male	0	
12	Other-service	Not-in-family	White	Female	0	
13	Exec-managerial	Not-in-family	White	Male	0	

	capital.loss	hours.per.week	native.country	income
7	3683	20	United-States	1
10	3004	35	United-States	1
11	2824	45	United-States	1
12	2824	20	United-States	1
13	2824	55	United-States	1

```
[ ]: workclass_types = df_more["workclass"].value_counts()
labels = list(workclass_types.index)
aggregate = list(workclass_types)
print(workclass_types)
print(aggregate)
print(labels)
```

```
Private          4876
Self-emp-not-inc  714
```

```

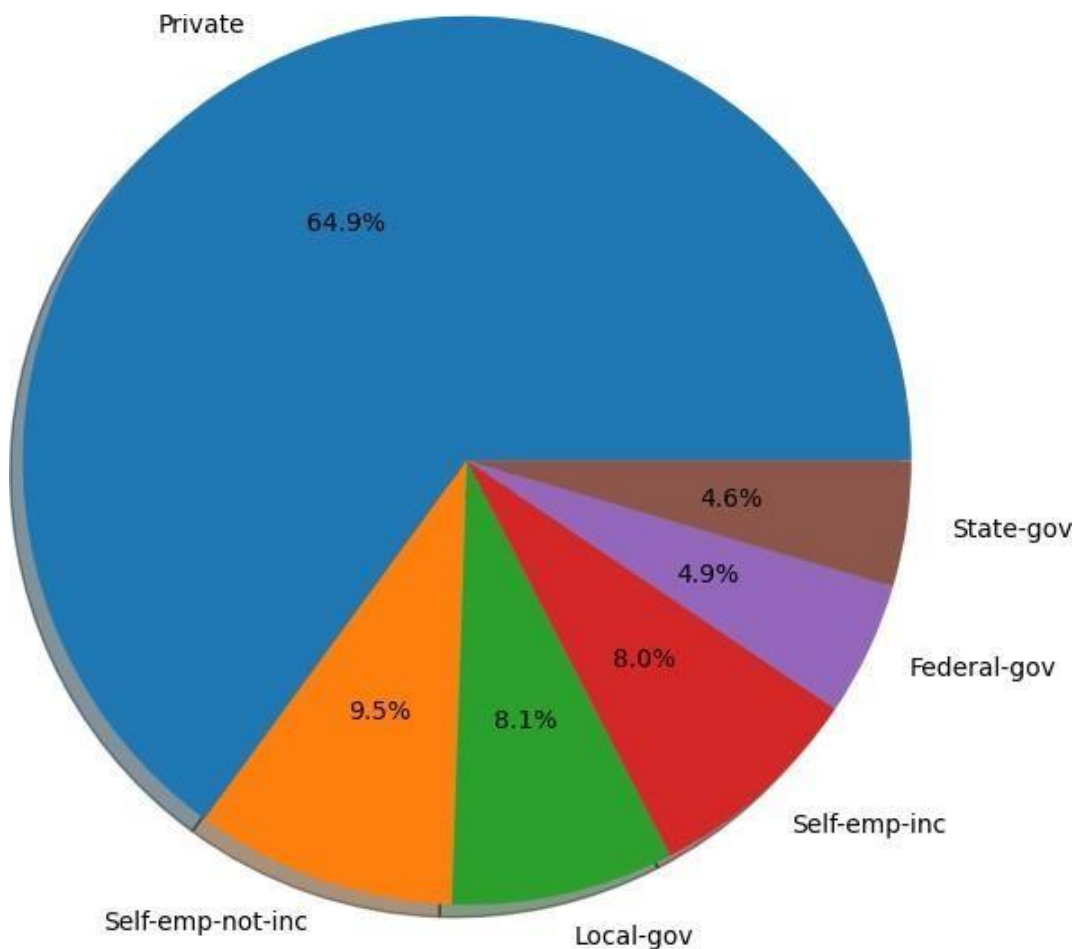
Local-gov          609
Self-emp-inc       600
Federal-gov        365
State-gov          344
Name: workclass, dtype: int64
[4876, 714, 609, 600, 365, 344]
['Private', 'Self-emp-not-inc', 'Local-gov', 'Self-emp-inc', 'Federal-gov',
'State-gov']

```

```

[ ]: plt.figure(figsize=(7,7))
plt.pie(aggregate, labels=labels, autopct='%1.1f%%', shadow = True)
plt.axis('equal')
plt.show()

```




```

[ ]: #Count plot on single categorical variable
sns.countplot(x = 'income', data = df)

```

```
plt.show()  
df["income"].value_counts()
```




```
[ ]: 0    22661  
     1     7508  
     Name: income, dtype: int64
```

```
[ ]: #Plot figsize  
     plt.figure(figsize=(10,7))  
     sns.heatmap(df.corr(), cmap="coolwarm", annot=True)  
     print(plt.show())
```

<ipython-input-91-6201d8194dba>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```



None

```
[ ]: plt.figure(figsize=(10,7))
sns.distplot(df["age"], color="red", bins=100)
plt.ylabel("Distribution", fontsize = 10)
plt.xlabel("Age", fontsize = 10)
plt.show()
```


<ipython-input-92-1b72b8b67fa9>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df["age"], color="red", bins=100)
```



[]: *#To find distribution of categorical columns w.r.t income*

```
fig, axes = plt.subplots(figsize=(20, 10))
```

```
plt.subplot(231)
sns.countplot(x = 'workclass',
              hue='income',
              data = df,
              palette="BuPu")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(232)
sns.countplot(x = 'marital.status',
              hue='income',
              data = df,
              palette="deep")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(233)
sns.countplot(x = 'education',
              hue='income',
              data = df,
              palette = "autumn")
```

```

plt.xticks(rotation=90)

plt.subplot(234)
sns.countplot(x = 'relationship',
              hue='income',
              data = df,
              palette = "inferno")
plt.xticks(rotation=90)

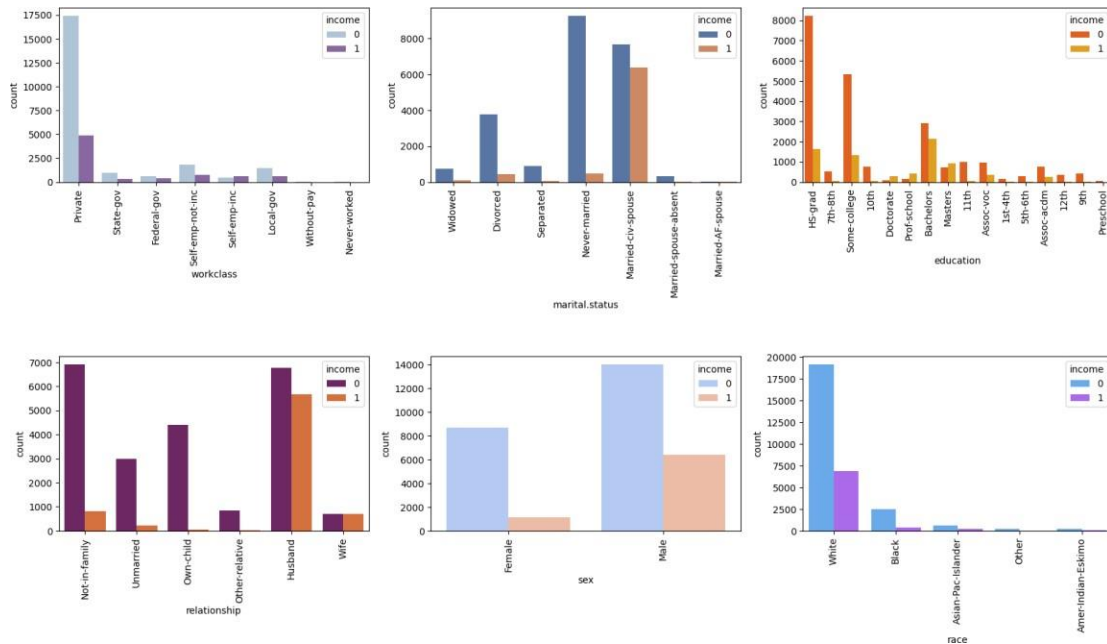
plt.subplot(235)
sns.countplot(x = 'sex',
              hue='income',
              data = df,
              palette = "coolwarm")
plt.xticks(rotation=90)

plt.subplot(236)
sns.countplot(x = 'race',
              hue='income',
              data = df,
              palette = "cool")
plt.xticks(rotation=90)
plt.subplots_adjust(hspace=1)
plt.show()

```

<ipython-input-93-42defcd4889b>:4: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call `ax.remove()` as needed.

```
plt.subplot(231)
```



```
[ ]: df1 = df.copy()
```

```
[ ]: categorical_features = list(df1.select_dtypes(include=['object']).columns)
print(categorical_features)
df1
```

```
['workclass', 'education', 'marital.status', 'occupation', 'relationship',
'race', 'sex', 'native.country']
```

```
[ ]:
    age workclass  fnlwgt  education  education.num  marital.status \
1      82   Private 132870    HS-grad             9      Widowed
3      54   Private 140359    7th-8th             4      Divorced
4      41   Private 264663  Some-college          10      Separated
5      34   Private 216864    HS-grad             9      Divorced
6      38   Private 150601    10th              6      Separated
...
32556  22   Private 310152  Some-college          10      Never-married
32557  27   Private 257302  Assoc-acdm          12  Married-civ-spouse
32558  40   Private 154374    HS-grad             9  Married-civ-spouse
32559  58   Private 151910    HS-grad             9      Widowed
32560  22   Private 201490    HS-grad             9      Never-married

    occupation  relationship  race  sex  capital.gain \
1  Exec-managerial  Not-in-family  White  Female      0
3  Machine-op-inspct  Unmarried  White  Female      0
4  Prof-specialty  Own-child  White  Female      0
```

5	Other-service	Unmarried	White	Female	0
6	Adm-clerical	Unmarried	White	Male	0
...
32556	Protective-serv	Not-in-family	White	Male	0
32557	Tech-support	Wife	White	Female	0
32558	Machine-op-inspct	Husband	White	Male	0
32559	Adm-clerical	Unmarried	White	Female	0
32560	Adm-clerical	Own-child	White	Male	0

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	0
3	3900	40	United-States	0
4	3900	40	United-States	0
5	3770	45	United-States	0
6	3770	40	United-States	0
...
32556	0	40	United-States	0
32557	0	38	United-States	0
32558	0	40	United-States	1
32559	0	40	United-States	0
32560	0	20	United-States	0

[30169 rows x 15 columns]

```
[ ]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for feat in categorical_features:
    df1[feat] = le.fit_transform(df1[feat].astype(str))
df1
```

```
[ ]: age workclass fnlwgt education education.num marital.status \
1 82 3 132870 11 9 6
3 54 3 140359 5 4 0
4 41 3 264663 15 10 5
5 34 3 216864 11 9 0
6 38 3 150601 0 6 5
...
32556 22 3 310152 15 10 4
32557 27 3 257302 7 12 2
32558 40 3 154374 11 9 2
32559 58 3 151910 11 9 6
32560 22 3 201490 11 9 4

occupation relationship race sex capital.gain capital.loss \
1 4 1 4 0 0 4356
3 7 4 4 0 0 3900
4 10 3 4 0 0 3900
```


5	8	4	4	0	0	3770
6	1	4	4	1	0	3770
...
32556	11	1	4	1	0	0
32557	13	5	4	0	0	0
32558	7	0	4	1	0	0
32559	1	4	4	0	0	0
32560	1	3	4	1	0	0

	hours.per.week	native.country	income
1	18	38	0
3	40	38	0
4	40	38	0
5	45	38	0
6	40	38	0
...
32556	40	38	0
32557	38	38	0
32558	40	38	1
32559	40	38	0
32560	20	38	0

[30169 rows x 15 columns]

```
[ ]: X = df1.drop(columns = ["income"])
y = df1["income"].values

# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
↳ 3, random_state = 0)

print ("Train set size: ", X_train.shape)
print ("Test set size: ", X_test.shape)
```

Train set size: (21118, 14)

Test set size: (9051, 14)

```
[ ]: from sklearn.ensemble import AdaBoostClassifier

# Train Adaboost Classifier
abc = AdaBoostClassifier(n_estimators = 300, learning_rate=1)
abc_model = abc.fit(X_train, y_train)

#Prediction
y_pred_abc = abc_model.predict(X_test)
```

```
print("Accuracy: ", accuracy_score(y_test, y_pred_abc))
print("F1 score :",f1_score(y_test, y_pred_abc, average="binary"))
print("Precision : ", precision_score(y_test, y_pred_abc))
```


Accuracy: 0.8637719588995691

F1 score : 0.7008007765105557

Precision :0.7921009325287987

```
[ ]: cm = confusion_matrix(y_test, y_pred_abc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap =_
    ↪"coolwarm");
plt.ylabel("Actual label");
plt.xlabel("Predicted label");
plt.title("Confusion Matrix - score:" + str(round(accuracy_score(y_test,_
    ↪y_pred_abc), 2)), size = 15);
plt.show()

print("confusion matrix\n",confusion_matrix(y_test,y_pred_abc))
print(classification_report(y_test, y_pred_abc))
```





confusion matrix

[[6374 379]

[854 1444]]

	precision	recall	f1-score	support
0	0.88	0.94	0.91	6753
1	0.79	0.63	0.70	2298
accuracy		0.86		9051
macro avg	0.84	0.79	0.81	9051
weighted avg	0.86	0.86	0.86	9051

